

Embedded System Software Design

Project 2

Problem Definition:

Since multi-core architectures are widely developed in most of hardware platforms, feasible task management mechanisms for multi-core systems are also proposed. In this project to input a real user application, and runs on a multi-core system. By parallelizing some regions in a program, we can reduce the execution duration, but make higher scheduling complexity. We compare FIFO, Round-Robin and CFS scheduler performance and temperature management in Linux system.

Experimental Environment:

- ✓ PC: i7 (8 cores) or 4 cores
- ✓ OS: Ubuntu 15.14
- ✓ Compiler: GCC 5.2.1

Scheduler setting

Linux supports several schedulers, such as FIFO 、Round-Robin and CFS. We can use the function “`sched_setscheduler(pid_t pid, int policy, const struct sched_param *param)`” to set the scheduling policy of the process specified by `pid` to `policy` and the scheduling parameters to “`param`”.

If `pid` is 0, the policy and parameters are set for the calling process. The following policies are available:

- SCHED_FIFO

First in first out. Processes are allowed on the CPU in the order in which they were added to the queue of processes to be run, for each priority.

- SCHED_RR

Round-Robin. Identical to SCHED_FIFO except that a process runs only for a set time slice (see `sched_rr_get_interval()`). Once the process has completed its time slice it is placed on the tail of the queue of processes to be run, for its priority.

- SCHED_OTHER

Non-realtime scheduling. This uses the traditional UNIX scheduler.

```
#include <sched.h>
```

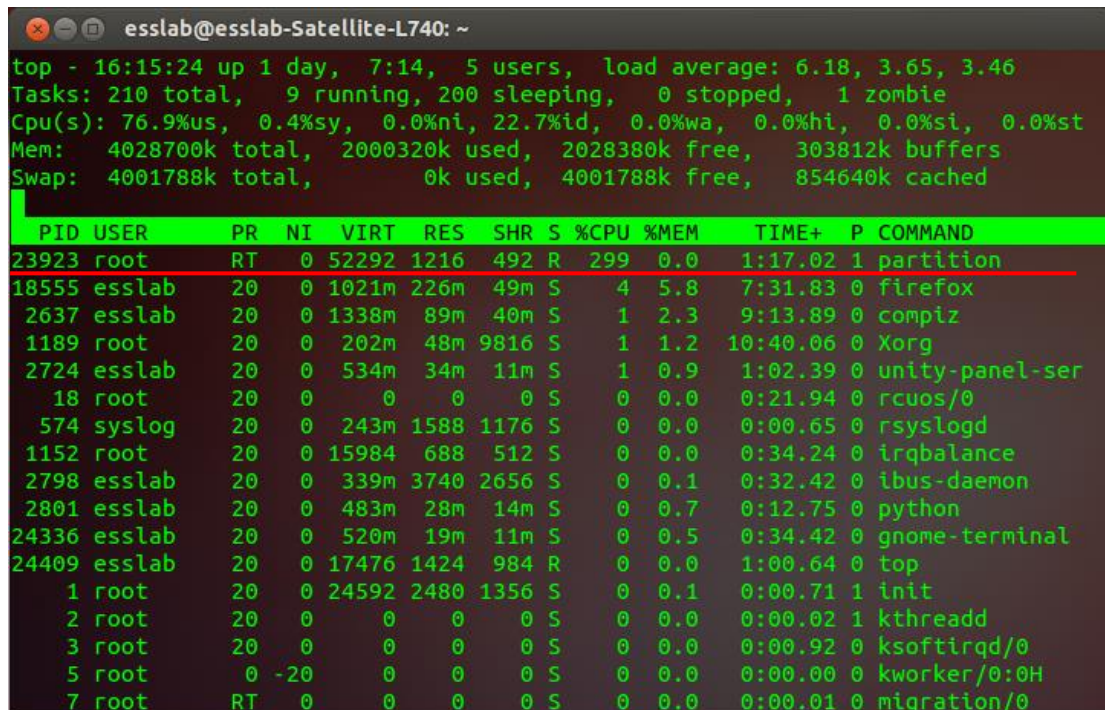
```
struct sched_param sp;
```

```
sp.sched_priority = sched_get_priority_max(SCHED_FIFO);
```

```
ret = sched_setscheduler(0, SCHED_FIFO, &sp);
```

In real-time scheduling policies, we need to set priority of the process. Linux allows the **static priority** value range 1 to 99 for SCHED_FIFO and SCHED_RR and the priority 0 for SCHED_OTHER (non-realtime). We can use the function “sched_get_priority_max” to get the range of the policies.

In the result, we can figure out that the priority of the process has been “RT” or negative values. It represents the process is running in real time policies and the priority is higher than normal processes.



```
top - 16:15:24 up 1 day, 7:14, 5 users, load average: 6.18, 3.65, 3.46
Tasks: 210 total, 9 running, 200 sleeping, 0 stopped, 1 zombie
Cpu(s): 76.9%us, 0.4%sy, 0.0%ni, 22.7%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Mem: 4028700k total, 2000320k used, 2028380k free, 303812k buffers
Swap: 4001788k total, 0k used, 4001788k free, 854640k cached
```

| PID | USER | PR | NI | VIRT | RES | SHR | S | %CPU | %MEM | TIME+ | P | COMMAND |
|-------|--------|----|-----|-------|------|------|---|------|------|----------|---|-----------------|
| 23923 | root | RT | 0 | 52292 | 1216 | 492 | R | 299 | 0.0 | 1:17.02 | 1 | partition |
| 18555 | esslab | 20 | 0 | 1021m | 226m | 49m | S | 4 | 5.8 | 7:31.83 | 0 | firefox |
| 2637 | esslab | 20 | 0 | 1338m | 89m | 40m | S | 1 | 2.3 | 9:13.89 | 0 | compiz |
| 1189 | root | 20 | 0 | 202m | 48m | 9816 | S | 1 | 1.2 | 10:40.06 | 0 | Xorg |
| 2724 | esslab | 20 | 0 | 534m | 34m | 11m | S | 1 | 0.9 | 1:02.39 | 0 | unity-panel-ser |
| 18 | root | 20 | 0 | 0 | 0 | 0 | S | 0 | 0.0 | 0:21.94 | 0 | rcuos/0 |
| 574 | syslog | 20 | 0 | 243m | 1588 | 1176 | S | 0 | 0.0 | 0:00.65 | 0 | rsyslogd |
| 1152 | root | 20 | 0 | 15984 | 688 | 512 | S | 0 | 0.0 | 0:34.24 | 0 | irqbalance |
| 2798 | esslab | 20 | 0 | 339m | 3740 | 2656 | S | 0 | 0.1 | 0:32.42 | 0 | ibus-daemon |
| 2801 | esslab | 20 | 0 | 483m | 28m | 14m | S | 0 | 0.7 | 0:12.75 | 0 | python |
| 24336 | esslab | 20 | 0 | 520m | 19m | 11m | S | 0 | 0.5 | 0:34.42 | 0 | gnome-terminal |
| 24409 | esslab | 20 | 0 | 17476 | 1424 | 984 | R | 0 | 0.0 | 1:00.64 | 0 | top |
| 1 | root | 20 | 0 | 24592 | 2480 | 1356 | S | 0 | 0.1 | 0:00.71 | 1 | init |
| 2 | root | 20 | 0 | 0 | 0 | 0 | S | 0 | 0.0 | 0:00.02 | 1 | kthreadd |
| 3 | root | 20 | 0 | 0 | 0 | 0 | S | 0 | 0.0 | 0:00.92 | 0 | ksoftirqd/0 |
| 5 | root | 0 | -20 | 0 | 0 | 0 | S | 0 | 0.0 | 0:00.00 | 0 | kworker/0:0H |
| 7 | root | RT | 0 | 0 | 0 | 0 | S | 0 | 0.0 | 0:00.01 | 0 | migration/0 |

Temperature Management & CPU Frequencies Scaling:

Running the applications will generate thermal for CPU, and if the CPU frequencies are getting higher the temperature will also be higher. We can use the thermal sensors in the CPUs to read the temperature value, and scale the CPU Frequencies to manage the temperature.

Use this command “#sudo apt-get install lm-sensors” to install the package and use command “#watch -n 0 sensors” to sense the CPU’s voltage & temperature

```
root@ubuntu1564bit-P5KPL-AM-BM: /home/ubuntu1564bit
Every 0.1s: sensors Thu Feb 9 04:10:57 2017

atk0110-acpi-0
Adapter: ACPI interface
Vcore Voltage:      +1.18 V (min = +0.85 V, max = +1.60 V)
+3.3 Voltage:      +3.26 V (min = +2.97 V, max = +3.63 V)
+5 Voltage:        +5.07 V (min = +4.50 V, max = +5.50 V)
+12 Voltage:       +12.14 V (min = +10.20 V, max = +13.80 V)
CPU FAN Speed:     1318 RPM (min = 600 RPM, max = 7200 RPM)
CHASSIS FAN Speed: 1074 RPM (min = 800 RPM, max = 7200 RPM)
POWER FAN Speed:   0 RPM (min = 800 RPM, max = 7200 RPM)
CPU Temperature:   +42.0°C (high = +60.0°C, crit = +95.0°C)
MB Temperature:    +39.0°C (high = +60.0°C, crit = +95.0°C)

coretemp-isa-0000
Adapter: ISA adapter
Core 0:            +43.0°C (high = +74.0°C, crit = +100.0°C)
Core 1:            +44.0°C (high = +74.0°C, crit = +100.0°C)
Core 2:            +47.0°C (high = +74.0°C, crit = +100.0°C)
Core 3:            +51.0°C (high = +74.0°C, crit = +100.0°C)
```

Check the **scaling_available_frequencies** file to know how many frequency levels that your CPU can be scaled.

Change the **scaling_governor** file to change the governor which being used right now.

Change the **scaling_setspeed** file to change frequency for current CPUs.

Change the governor to “**userspace**” that we can scale the frequency by ourselves.

Using “cat” to print the information in the file.

```
cat /XXX/XXX/.../cpu0/....frequency
cat /XXX/XXX/.../cpu1/....frequency
cat /XXX/XXX/.../cpu2/....frequency
cat /XXX/XXX/.../cpu3/....frequency

cat /XXX/XXX/.../cpu0/....governor
cat /XXX/XXX/.../cpu1/....governor
cat /XXX/XXX/.../cpu2/....governor
cat /XXX/XXX/.../cpu3/....governor
exit 0
```

Using “echo” to change the value in the file

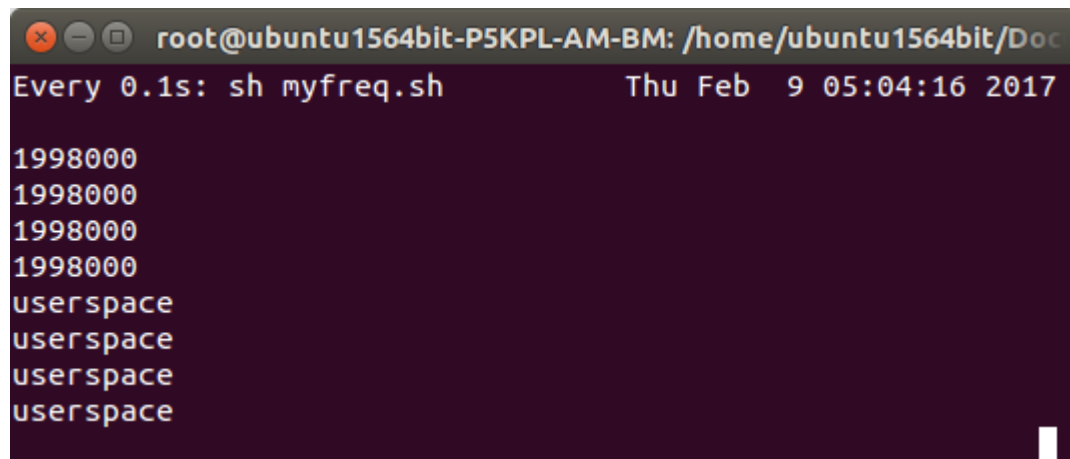
```

echo ??? >> /XXX/XXX/.../cpu0/....frequency
echo ??? >> /XXX/XXX/.../cpu1/....frequency
echo ??? >> /XXX/XXX/.../cpu2/....frequency
echo ??? >> /XXX/XXX/.../cpu3/....frequency

echo ??? >> /XXX/XXX/.../cpu0/....governor
echo ??? >> /XXX/XXX/.../cpu1/....governor
echo ??? >> /XXX/XXX/.../cpu2/....governor
echo ??? >> /XXX/XXX/.../cpu3/....governor
exit 0

```

Use `#watch -n 0 sh XXXX.sh` to run the scripts continuously.



```

root@ubuntu1564bit-P5KPL-AM-BM: /home/ubuntu1564bit/Doc
Every 0.1s: sh myfreq.sh Thu Feb 9 05:04:16 2017

1998000
1998000
1998000
1998000
userspace
userspace
userspace
userspace

```

Low Frequency Setting:

```

root@ubuntu1564bit-P5KPL-AM-BM:/home/ubuntu1564bit/Documents/project2/final/10#
time ./partition10_default

real    0m47.439s
user    2m19.192s
sys     0m0.004s
root@ubuntu1564bit-P5KPL-AM-BM:/home/ubuntu1564bit/Documents/project2/final/10#
time ./partition10_FIFO

real    0m44.735s
user    2m13.152s
sys     0m0.004s
root@ubuntu1564bit-P5KPL-AM-BM:/home/ubuntu1564bit/Documents/project2/final/10#
time ./partition10_RR

real    0m45.245s
user    2m12.728s
sys     0m0.000s
root@ubuntu1564bit-P5KPL-AM-BM:/home/ubuntu1564bit/Documents/project2/final/10#

```

| | | | |
|---------|---------|---------|---------|
| Core 0: | +48.0°C | Core 0: | +57.0°C |
| Core 1: | +49.0°C | Core 1: | +58.0°C |
| Core 2: | +52.0°C | Core 2: | +62.0°C |
| Core 3: | +57.0°C | Core 3: | +66.0°C |

High Frequency Setting:

```
root@ubuntu1564bit-P5KPL-AM-BM: /home/ubuntu1564bit/Documents/project2/final/10#  
time ./partition10_default  
  
real    0m43.543s  
user    2m9.104s  
sys     0m0.008s  
root@ubuntu1564bit-P5KPL-AM-BM: /home/ubuntu1564bit/Documents/project2/final/10#  
time ./partition10_FIFO  
  
real    0m42.514s  
user    2m4.352s  
sys     0m0.000s  
root@ubuntu1564bit-P5KPL-AM-BM: /home/ubuntu1564bit/Documents/project2/final/10#  
time ./partition10_RR  
  
real    0m42.465s  
user    2m5.912s  
sys     0m0.000s  
root@ubuntu1564bit-P5KPL-AM-BM: /home/ubuntu1564bit/Documents/project2/final/10#
```

| | | | |
|---------|---------|---------|---------|
| Core 0: | +51.0°C | Core 0: | +58.0°C |
| Core 1: | +51.0°C | Core 1: | +59.0°C |
| Core 2: | +56.0°C | Core 2: | +65.0°C |
| Core 3: | +59.0°C | Core 3: | +72.0°C |

Crediting :

The number of created threads must be twice as the number of CPUs, and the number of CPUs should be more than 3.

[Scheduler Implementation. 40%]

- Describe how to implement the scheduler setting (FIFO, RR, Default) **20%**
- Show the scheduling states of tasks **20%**

[Temperature Watch & Frequency Scaling. 40%]

- Show that your computer's frequency levels and write the scripts 20%
- Show that CPU's temperature in different frequency setting. 20%

[Result. 20%]

- Compare the response time of the program in three scheduler setting (FIFO, RR, Default) **10%**
- Analyze the performance of three scheduler setting (FIFO, RR, Default) and in different frequency settings. **10%**

Project submit

Submit deadline: 09 :00, May. 12, 2017

Submission : [Moodle](#)

File name format : ESSD_Student ID_HW2.rar

※ Strictly prohibited copying !

ESSD _Student ID_HW2.rar must include the **report** and **source code**.

嚴禁抄襲，發生該類似情況者，一律以零分計算